

Smoke Benchmark Report

- Campaign: 20260713-135148-subscription-campaign
- Generated: 2026-07-13T14:22:20.2961160+02:00
- Verdict: **PASS**
- Comparative conclusion: **INCONCLUSIVE**

Executive result

All 4/4 approved calls passed public tests, hidden grading at 100/100 and applicable enforcement. There were 0 overlapping executions and 0 invalid attempts. This proves smoke viability of the repaired harness; it does not prove that the full arm outperforms vanilla.

Run evidence

Provider	Arm	Model evidence	Seconds	Public	Hidden	Enforcement	Protected changes
claude	core-router-enforcement	claude-opus-4-8 (provider-event)	90.1	True	100	True	0
codex	vanilla	unresolved-provider-default (not-reported)	59	True	100	True	1
codex	core-router-enforcement	unresolved-provider-default (not-reported)	93.3	True	100	True	0
claude	vanilla	claude-opus-4-8 (provider-event)	69.3	True	100	True	0

Integrity and review

- Harness snapshot: 5 pinned file hashes.
- Full-arm protected-input integrity: True.
- Total provider wall time: 311.7 seconds.
- Immediately after stage completion, all five pinned harness hashes matched and exactly four provider executions existed since campaign creation.
- Codex vanilla added useful regression cases inside the existing public test file; it did not delete or weaken assertions, but this is reported as a protected-input change that the full arm avoided by creating a separate regression file.

Publication blockers

- Only one trial per provider/arm cell; comparative effect is statistically inconclusive.
- Codex provider-default did not report an exact resolved model; post-smoke stages require explicit model selectors.

Next gate

Directional remains unapproved. Create a new campaign with explicit model selectors before any post-smoke calls, then require separate human approval.

Structured JSON beside this report is the measurement source of truth. The PDF is a human-readable artifact.