

Calibration Benchmark Report

- Campaign: 20260713-151556-calibration-probe
- Generated: 2026-07-13T16:00:02.1886024+02:00
- Verdict: **PASS**
- Comparative conclusion: **CEILING**

Executive result

All 2/2 approved calls passed public tests, hidden grading at 100/100 and applicable enforcement. There were 0 overlapping executions and 0 invalid attempts. Both arms reached the ceiling, so this fixture cannot discriminate between vanilla and the full environment. The predeclared decision is STOP: spend zero Claude calls and redesign a harder fixture.

Run evidence

Provider	Arm	Model evidence	Seconds	Public	Hidden	Enforcement	Protected changes
codex	vanilla	gpt-5.6-sol (explicit-cli-selector)	77.5	True	100	True	0
codex	core-router-enforcement	gpt-5.6-sol (explicit-cli-selector)	138.7	True	100	True	0

Integrity and review

- Harness snapshot: 5 pinned file hashes.
- Full-arm protected-input integrity: True.
- Total provider wall time: 216.2 seconds.
- Both implementations changed only src/migration.py and are semantically equivalent: schema introspection, idempotent upgrade, reversible downgrade, row and UNIQUE preservation. Vanilla provider wall time was 77.5 seconds; full-arm provider wall time was 138.7 seconds; this two-call screening probe is not a latency estimate.

Publication blockers

- Screening-only calibration is excluded from publishable confirmatory scores.

Next gate

STOP: both arms scored 100. Spend zero Claude calls and redesign a harder fixture before requesting any new provider approval.

Structured JSON beside this report is the measurement source of truth. The PDF is a human-readable artifact.