Agentic Work Best Practices - Architecture Transformation

Stable Core, Objective Integrity, Router, evals, and candidate-only research

2026-07-12

Source of truth · SebsKk/AgenticWorkBestPractices

Outcome

The Claude v1 prototype was preserved as a ZIP with a SHA-256 checksum. Its research and long contracts remain available as reference, but they are no longer loaded as universal law. Active bootstraps now point to a small promoted Stable Core and task-specific Context Packs.

What changed

- Added an immutable objective record and a machine-checkable ledger with 29 requirements across eight pillars.
- Added `objective-check`, which refuses completion while requirements lack evidence.
- Added a deterministic Agentic Knowledge Router with explicit selected/rejected modules and context budgets.
- Added `init` and `doctor` tooling for idempotent pointer management and Windows/OneDrive diagnostics.
- Added twelve mock-case definitions and four controlled comparison arms: vanilla, Core, Core+Router, and Core+Router+enforcement.
- Ran the deterministic routing suite: **12/12 cases passed**. Full agent outcome trials remain intentionally unclaimed until provider adapters and hidden fixtures are isolated.
- Replaced direct research-to-stable mutation with candidate-only research and separate promotion.
- Added a durable registry of 21 sources with evidence tier, topics, status, review date and next recheck date.
- Added official, academic, benchmark, security, YouTube, podcast and practitioner/social research lanes.
- Replaced blanket test blocking and generic `npm test` hooks with candidate templates protecting only immutable eval material.
- Switched active global and sibling-repository pointers to the minimal Stable manifest while preserving local project rules.

Impact

Agents no longer need to consume the complete knowledge base at startup. Parser/discovery work receives Generalization Assurance; API/schema changes receive Change Impact; production/security work receives security boundaries; research receives candidate governance. The objective ledger exposes omitted requirements instead of allowing optimistic self-completion.

Current limitations

- The routing benchmark validates module selection, not final coding-agent quality.

- Provider-specific outcome trials require at least five trials per case and arm with identical models, budgets and hidden graders.
- The long v1 contracts contain claims that still require claim-level migration and evaluation before promotion.
- Hook templates remain candidates until tested against installed Claude/Codex versions in representative repositories.

Sources

- OpenAI Codex documentation - tier 1, checked 2026-07-12.
- OpenAI Codex changelog - tier 1, checked 2026-07-12.
- Claude Code documentation - tier 1, checked 2026-07-12.
- Claude Code changelog - tier 1, checked 2026-07-12.
- Anthropic Engineering - tier 1, checked 2026-07-12.
- Evaluating AGENTS.md - primary paper, checked 2026-07-12.
- Probe-and-Refine Repository Guidance - primary paper, checked 2026-07-12.
- ContextBench - primary paper, checked 2026-07-12.
- METR developer productivity study - primary study, checked 2026-07-12.
- SWE-bench - benchmark, checked 2026-07-12.
- OpenHands Index - benchmark, checked 2026-07-12.
- OWASP MCP Security - security guidance, checked 2026-07-12.
- NIST SP 800-218A - standard, checked 2026-07-12.
- Anthropic YouTube - discovery lane, checked 2026-07-12.
- OpenAI YouTube - discovery lane, checked 2026-07-12.
- AI Engineer talks - discovery lane, checked 2026-07-12.
- Latent Space podcast - discovery lane, checked 2026-07-12.
- Simon Willison AI notes - practitioner signal, checked 2026-07-12.
- Geoffrey Huntley - Ralph - practitioner signal, checked 2026-07-12.
- Addy Osmani - practitioner signal, checked 2026-07-12.