

Benchmarking Quick Reference

1. Pass infrastructure and `AgenticBench` doctor gates.
2. Confirm both subscription logins and a provider snapshot younger than seven days.
3. Create the campaign without starting a model.
4. Approve only the next stage: smoke, directional, confidence or pilot.
5. Use one terminal and one runner only; never start a second runner to recover a silent window.
6. Execute only inside the copied public workspace.
7. Grade public behavior and the hidden case on the host after agent exit.
8. Stop on `Evals/local/STOP`, timeout, auth drift, isolation drift, orphan process or two consecutive failures.
9. Complete five trials per fixture/provider/arm cell only after earlier gates justify the spend.
10. Summarize uncertainty and promote rules separately with human approval and rollback.

Pilot: `2 fixtures x 2 providers x 2 arms x 5 trials = 40 runs`.

Spend ladder: `4 -> 8 -> 24 -> 40` cumulative runs. Start with four, not forty.

Hold fixed: resolved model, effort, prompt, fixture, permissions, budgets, grader and environment.

Never claim that routing equals coding quality, public tests equal final-state proof, logical separation equals hidden isolation, an alias equals the resolved model, SWE-bench Verified is the primary 2026 frontier metric, or a benchmark score automatically measures human productivity.

Smoke may use `provider-default`; every later stage requires explicit models. Publication requires complete cells, inaccessible grader, resolved model telemetry, approved stage budgets, no concurrency, no leakage, retained failures and reported uncertainty.