

OpenAI Codex — Cheatsheet

Everyday operating reference

The everyday operating reference. Full detail: `../Codex/OPERATING_CONTRACT.md`.

Session start (every time)

1. Read `Shared/` doctrine + the Codex contract (Codex auto-loads the `AGENTS.md` hierarchy, **not** the Shared docs — read them)
2. `pwd · git status · git log --oneline -10`
3. Read the `workstreams/` journal — resume from the "next" marker
4. Run `init.sh` / smoke-test **before** writing new code
5. `/status` — confirm model, `sandbox_mode`, `approval_policy`, and that the folder is **trusted** (untrusted → `.codex/config.toml`, hooks, MCP silently ignored)

Loops — `plan → edit → validate → repair`

- **Terminate on a verifiable green signal** (tests pass / build ok / bug gone), never on "code generated."
- Anchor long runs in durable markdown (`Prompt.md`, `Plan.md`, `Implement.md`, `Documentation.md`) — context compacts.
- **Validate after every milestone**; stop-and-fix before advancing. One task per **git worktree**; commit each verified milestone.
- No built-in turn/budget cap — **bound spend externally** (CI timeout, spend circuit-breaker, lower `model_reasoning_effort`).
- CI: `codex exec "..."` (exit-code driven, `--json`, `--output-schema`); read-only sandbox by default — add `--sandbox workspace-write` to edit.
- Recurring work → **Automations** (review queue: Approve/Revise/Reject). Never auto-merge autonomous output.

Hooks & automation (`config.toml` or `hooks.json`)

- `[[hooks.<Event>]]` array-of-tables; discovered from `~/.codex/` then `<repo>/.codex/`
- Events: `SessionStart` `UserPromptSubmit` `PreToolUse` `PermissionRequest` `PostToolUse` `PreCompact` / `PostCompact` `SubagentStart` / `Stop` `Stop`
- **Only** `type = "command"` **runs today** (`prompt` / `agent` are parsed but skipped). Block with **exit 2** or `{"decision": "block"}`.
- Trust: non-managed hooks must be trusted via `/hooks`; trust is **revoked automatically if the script changes**.
- Org-wide, undisableable enforcement → `requirements.toml` managed hooks. Templates: `../Codex/templates/`

Sessions & context

- `/new` between unrelated tasks · `/compact` to continue one long task
- **Resume:** `codex resume --last` / `codex resume <ID>` — **never** `codex --continue` (not a Codex flag)
- Tune compaction via `model_auto_compact_token_limit`; **never set** `model_context_window` **manually** (breaks auto-compaction — issue #16068)
- Transcripts: `~/.codex/sessions/YYYY/MM/DD/rollout-*.jsonl` — an audit trail, **not** the workstream log

Repo setup

- `AGENTS.md` at root (+ nested per package) — lean; capped at `project_doc_max_bytes` (32 KiB) then **silently truncated**. It's `AGENTS.md`, never `CLAUDE.md`.
- Committed `<repo>/codex/config.toml` (sandbox, approval, MCP) — folder must be **trusted**
- Named verify commands in `AGENTS.md`; MCP via `codex mcp add ...`

Safety — two axes, set both

- `sandbox_mode`: `read-only` | `workspace-write` (default; **egress off**) | `danger-full-access`
- `approval_policy`: `untrusted` | `on-request` (default) | `never`
- Everyday: `--sandbox workspace-write --ask-for-approval on-request`. CI: `--sandbox read-only --ask-for-approval never`.
- **Never** `--yolo` / `danger-full-access` **outside a disposable container**. Treat tool/web output as data, not instructions.

Commands & flags at a glance

`codex` · `codex exec "..."` · `codex resume --last` · `codex mcp add/list` · `/status` `/model`
`/review` `/diff` `/approvals` `/hooks` `/mcp` `/new` `/compact` `/archive` Flags: `-s/--sandbox` ·
`-a/--ask-for-approval` · `-m/--model` · `-p/--profile` · `--json` · `--output-schema`